

Google Cloud Storage (GCS)

1.Introduction

What is GCS?

Google Cloud Storage (GCS) is an object storage service commonly used as the storage foundation for data lakes. It stores unstructured, semi-structured, or structured files such as raw text, PDFs, images, Parquet, JSON, and CSVs in containers called Buckets.'

Why is GCS Used?

- Cheap, Scalable File Storage: Storing raw files directly inside a database or analytics engine is expensive. GCS provides extremely cheap, virtually infinite storage.
- Separation of Storage from Compute: It lets you hold raw data at rest without paying for processing until you are ready to query or process it.

When to Use GCS?

- Data Landing Zone: Storing incoming raw files from external APIs, webhooks, or database exports before cleaning them.
- Storing Unstructured Media: Holding non-relational files like images, video clips, audio logs, or PDFs.
- Data Archiving & Backups: Storing historical logs or compliance backups using low-cost storage tiers such as Nearline, Coldline, and Archive.
- Intermediate Staging for Pipelines: Temporarily staging data while ETL pipelines such as Dataflow or Spark process it.

Name	Created	Location type	Location	Default storage class	Last modified
asia-south1-pub-sub-data-0834c28d...	Aug 19, 2026, 11:59:41 AM	Region	asia-south1	Standard	Aug 19, 2026, 4:00:15 PM
avrobucket23	Jun 16, 2026, 2:50:28 PM	Region	asia-south1	Standard	Jun 16, 2026, 2:50:28 PM
bronzelayer2	Jul 16, 2026, 12:50:58 PM	Region	asia-south1	Standard	Jul 16, 2026, 12:50:58 PM
cloudfunction23	Jun 23, 2026, 6:44:19 PM	Region	asia-south1	Standard	Jun 24, 2026, 11:55:46 AM
customers145	Jun 5, 2026, 3:25:33 PM	Region	asia-south1	Standard	Aug 26, 2026, 10:35:57 AM
dagstobq	Jun 18, 2026, 4:18:56 PM	Region	asia-south1	Standard	Jun 18, 2026, 4:18:56 PM
datacloud2	Jun 12, 2026, 3:58:14 PM	Region	asia-south1	Standard	Jun 12, 2026, 3:58:14 PM
dataflowbucket34	Jun 11, 2026, 3:44:16 PM	Region	asia-south1	Standard	Jun 11, 2026, 3:44:16 PM
dataflowcloushell	Jun 12, 2026, 1:34:24 PM	Region	asia-south1	Standard	Jun 12, 2026, 1:34:24 PM
dataflowscheduling	Jun 16, 2026, 11:32:58 AM	Region	asia-south1	Standard	Jun 16, 2026, 11:32:58 AM
datapipeline_bucket23	Jun 17, 2026, 1:02:57 PM	Region	asia-south1	Standard	Jun 17, 2026, 1:02:57 PM
dataproc-staging-asia-south1-221098...	Jul 1, 2026, 12:43:55 PM	Region	asia-south1	Standard	Jul 1, 2026, 12:43:55 PM
dataproc-temp-asia-south1-2210989...	Jul 1, 2026, 12:43:54 PM	Region	asia-south1	Standard	Jul 1, 2026, 12:43:54 PM
gcs-bq-pipeline-sa	Aug 17, 2026, 3:34:03 PM	Region	asia-south1	Standard	Aug 17, 2026, 3:34:03 PM
incoming-data-cf	Jul 21, 2026, 2:40:00 PM	Region	asia-south1	Standard	Jul 21, 2026, 2:49:50 PM
json_bucket123	Aug 5, 2026, 3:49:15 PM	Region	asia-south1	Standard	Aug 5, 2026, 3:49:15 PM
my-excel-archive	Jul 17, 2026, 11:31:17 AM	Region	asia-south1	Standard	Jul 17, 2026, 11:31:17 AM

2. GCS Core Concepts

Bucket

A Bucket is a container used to store objects in GCS.

A bucket provides the storage location where files and other objects are maintained.

Object

An Object is an individual file stored inside a GCS bucket.

Examples include:


- CSV files
- JSON files
- Parquet files
- PDF files
- Images
- Videos
- Logs

Storage Class

A Storage Class determines how data is stored based on how frequently it is expected to be accessed.

GCS provides different storage classes for frequently accessed, occasionally accessed, and rarely accessed data.

3.IAM Roles



IAM roles control **who can access your GCS buckets and objects**.

- **Storage Admin:** Full control over buckets and objects.
- **Storage Object Admin:** Can create, read, update, and delete objects, but cannot manage the bucket itself.
- **Storage Object User:** Can read, create, and delete objects.
- **Storage Object Viewer:** Can only **read** objects.
- **Storage Object Creator:** Can only **create** new objects. It cannot read or delete existing objects.

Uniform Bucket-Level Access (UBLA)

UBLA makes **IAM the main way to control access** to a bucket.

- It disables object-level ACLs.
- Permissions are managed at the **bucket level**.
- It helps keep permissions simple and consistent.

Public vs. Private Access

Private Access:

Only authorized users or applications can access the data. This is the **default**.

Public Access:

Anyone on the internet can access the objects when public permissions are granted.

Public Access Prevention (PAP):

Prevents a bucket from accidentally becoming publicly accessible.

4.GCS Bucket Locations

1. Region

Data is stored in **one geographic region** across multiple zones.

Example: asia-south1, us-central1

Use when: You want **lower cost and low latency** for applications running in the same region.

2. Dual-Region

Data is stored in **two different regions**.


Use when: You need **higher availability and protection from a regional failure**.

3. Multi-Region


Data is stored across **multiple regions in a large geographic area**.

Example: us, eu, asia

Use when: You need to serve users from **different locations around the world** and want high availability.



5. Encryption



GCS **encrypts your data automatically** when it is stored.

Key Management Models

1. Google-Managed Encryption Keys (Default)

- Google manages the encryption keys.
- No configuration is required.
- Google handles key rotation and management.
- No extra key management work is required.

2. Customer-Managed Encryption Keys (CMEK)

- You create and manage the keys using **Cloud KMS**.
- You have more control over the keys.
- Useful when your organization has strict security or compliance requirements.

3. Customer-Supplied Encryption Keys (CSEK)

- You provide your own encryption key when uploading or accessing data.
- GCS does not store the key.
- You are responsible for managing the key.

Why is Encryption Important?

- **Data Protection:** Protects your data if the physical storage hardware is accessed without authorization.
- **Compliance:** Helps organizations meet security and regulatory requirements.
- **Key Control:** CMEK gives you more control over encryption keys, including managing and rotating them.


Easy Way to Remember

Google-Managed Keys → Google manages the keys

CMEK → You manage keys using Cloud KMS

CSEK → You provide and manage your own keys

6.GCS Object Naming & Folder Structure



Google Cloud Storage (GCS) uses a **flat namespace**, meaning objects are stored in a single bucket without a traditional folder hierarchy.

Object Names and Prefixes

- **Object Name:** The complete name used to identify an object.
 - Example: telemetry/sensor_a/2026/08/26/metrics.parquet
- **Prefix:** The part of an object name before the actual file name.
 - Example: telemetry/sensor_a/2026/08/26/
- The / character is commonly used as a delimiter to create a **folder-like structure**.
- These folders are virtual; GCS does not create traditional directories.

Bucket Naming

Separate buckets based on their purpose or data lifecycle:


--acme-prod-raw-uscentral1

--acme-prod-curated-uscentral1

- **Raw:** Stores original/ingested data.

- **Curated:** Stores cleaned and transformed data

7.GCS Storage Classes



7.1 Standard Storage ("Hot Data")

What is it?

The default storage class designed for data that is accessed frequently, changed constantly, or read by live applications.

Why do we use it?

- **Zero Retrieval Fees:** You pay the highest rate per GB to store it, but you pay no extra fees whenever you read or download the data.
- **No Minimum Retention:** You can store a file for a short period, delete it, and only

pay for the time it was stored.

- **Maximum Speed & Availability:** High uptime SLA so your live applications are not blocked.

When to Use It?

Use it when files are read or updated multiple times a day, week, or month.

Real-World Examples

1. **Streaming App Assets:**

Storing movie thumbnail images or video clips for a web application like YouTube or Netflix so millions of users can stream them instantly.

2. **Daily ETL Ingestion:**

Storing incoming hourly raw CSV files that Dataflow pipelines actively read, clean, and load into BigQuery.

7.2 Nearline Storage ("Warm Data")

What is it?

A low-cost storage class designed for data you expect to access infrequently—typically once a month or less.

Why do we use it?

- **Lower Monthly Storage Cost:** Storing files here costs less per GB than Standard Storage.
- **The Trade-off:** You pay a small retrieval fee per GB whenever you read or download the data, and there is a 30-day minimum storage commitment.

When to Use It?

Use it for data you need to keep available but only access for monthly routines or occasional checks.

Real-World Examples

1. **Monthly Financial Reporting:**

Storing monthly sales spreadsheets and invoice PDFs that accounting opens once at the end of each month.

2. **App Crash Logs:**

Saving system logs from mobile applications. Engineers rarely look at them unless a bug is reported once every few weeks.

7.3 Coldline Storage ("Cold Data")

What is it?

A very low-cost storage class designed for data you expect to access at most once a quarter (every 90 days).

Why do we use it?

- **Significant Storage Discount:** Storing files here costs much less than Standard Storage.
- **The Trade-off:** Higher retrieval fees than Nearline, and a 90-day minimum storage commitment. If you delete the file before 90 days, you still pay for 90 days of storage.

When to Use It?

Use it for disaster recovery backups or quarterly auditing records.

Real-World Examples

1. **Disaster Recovery Backups:**

Storing quarterly database snapshots of a company's database so they can recover if a catastrophic outage occurs.

2. **Quarterly Compliance Reports:**

Storing tax calculation files that external auditors review only once every 3 months.

7.4 Archive Storage ("Frozen Data")

What is it?

The lowest-cost storage class designed for long-term data archiving, online backups, and regulatory compliance that you plan to access less than once a year.

Why do we use it?

- **Rock-Bottom Storage Cost:** Storing data costs a fraction of a cent per GB per month.
- **Instant Access:** GCS Archive data is accessible when you need it without a

traditional archive thawing process.

- **The Trade-off:** Very high retrieval fees and a 365-day (1 year) minimum storage commitment.

When to Use It?

Use it when legal regulations require you to hold onto data for years, but you expect it to be rarely accessed.

Real-World Examples

1. **Healthcare Medical Records:**

Storing patient X-ray images for several years to comply with healthcare requirements.

2. **Bank Account Transaction Archives:**

Keeping old closed bank account history in case it is required for a tax investigation.

7.5 Autoclass & Lifecycle Rules

As a Data Engineer, you don't want to manually move thousands of files between storage classes every month. GCP provides two solutions.

1. **GCS Lifecycle Rules**

You can set automated policies such as:

"Move files from Standard to Nearline after 30 days, then to Archive after 365 days."

2. **Autoclass**

Autoclass is a hands-off feature where Google automatically monitors how often a file is accessed and transitions it between classes to optimize storage costs without requiring you to write individual rules.


7.6 Storage Class Comparison

Storage Class	Intended Access Frequency	Min. Duration	Retrieval Fee	Best Analogy
Standard	Multiple times / month (Hot)	None	\$0	Your Everyday Wallet – Cash you use daily
Nearline	~1 time / month (Warm)	30 Days	Low	Your Piggy Bank – Open occasionally
Coldline	~1 time / quarter (Cold)	90 Days	Medium	A Home Safe – Opened for emergencies
Archive	< 1 time / year (Frozen)	365 Days	High	A Bank Vault – Locked away for years

8.GCS Object Lifecycle Management (OLM)

The screenshot shows the Google Cloud console interface for the 'gcs-bucket113' bucket. The 'Lifecycle' tab is selected, displaying a table with lifecycle rules. The table has three columns: 'Action', 'Object condition', and 'Works with'. A single rule is listed: 'Set to Coldline' with the condition '30+ days since object was created'. The console also includes a sidebar with navigation options like Overview, Buckets, Monitoring, Settings, and Storage Intelligence. The top of the console shows the project name 'My Project 90688' and a search bar.

Action	Object condition	Works with
Set to Coldline	30+ days since object was created	



Google Cloud Storage Object Lifecycle Management (OLM) is an automated rule engine configured at the bucket level. It evaluates objects asynchronously and applies predefined actions, such as storage class transitions or deletions, when specified conditions are satisfied.

Rule Logic

IF Condition THEN Action

8.1 What is OLM?

- A zero-code automation tool native to GCP Cloud Storage.
- Configured using JSON or through the Cloud Console / gcloud CLI.
- Governed by rule sets consisting of Actions and Conditions.

8.2 Why Do We Use OLM?

- **Cost Optimization:** Automatically moves idle or legacy data to lower-cost tiers such

as Nearline, Coldline, and Archive without manual oversight.

- **Storage Governance & Hygiene:** Cleans up intermediate pipeline data, failed multipart uploads, or temporary files automatically.
- **Compliance & Risk Mitigation:** Enforces data retention boundaries by removing records after their required lifecycle expires.

8.3 When to Use OLM?

- When managing large-scale, automated ETL/ELT pipelines.
- When enforcing company compliance or retention limits, such as "Delete logs after 90 days."
- When handling versioned buckets where legacy file versions accumulate and increase costs.


8.4 How Does OLM Operate?

- **Background Scanning:** GCS regularly inspects objects in a bucket against active rules.
- **Asynchronous Execution:** Once conditions are met, actions are scheduled and executed asynchronously by Google Cloud.

Supported Actions

Action	Description	Notes
SetStorageClass	Downgrades the storage tier, such as Standard → Nearline → Coldline → Archive.	One-way transition in the documented flow.
Delete	Deletes matching objects.	Subject to soft-delete and retention holds.
AbortIncompleteMultipartUpload	Cancels unfinished multipart file uploads.	Clears stuck upload chunks to save cost.

9. Object Versioning in GCS



Google Cloud Storage (GCS) Object Versioning is a bucket-level security and data integrity feature. When enabled, replacing or deleting a file does not permanently destroy it. Instead, GCS retains historical versions by assigning a unique generation number to each version of the object.

9.1 What is Object Versioning?

- By default, overwriting data.csv in GCS destroys the old version.
- With Object Versioning enabled, GCS retains historical versions of the file.

Live Version

The current, active version visible during normal file searches or API calls.

Noncurrent Versions

Archived historical snapshots of the file that remain accessible using their specific generation numbers.

9.2 Why Do We Use Object Versioning?

- Accidental Deletion & Overwrite Protection: Helps restore critical data if a bad

deployment, buggy script, or human mistake overwrites or deletes production files.

- **Auditability & Traceability:** Retains an immutable history of how data files evolved over time.
- **Seamless Disaster Recovery:** Allows rollback to a known good state.

9.3 When to Use Object Versioning?

- When storing critical configuration files, operational lookup tables, or master data that multiple teams update.
- When running pipeline tasks that frequently update destination files in place.
- When NOT to use it: On high-frequency temporary streaming staging buckets where thousands of temporary files are replaced frequently, because this can increase storage costs.

9.4 How Does It Operate?

Overwrite Behavior

Uploading a new file with the same name preserves the existing file as a noncurrent version and creates a new live version.

Delete Behavior

Deleting a file makes the file appear deleted while historical versions can remain available. Permanently removing a historical version requires deleting its specific generation.

10. GCS Data Recovery Concepts

10.1 Soft Delete – Your "Trash Bin"

What is it?

A safety feature that holds deleted files or buckets in a hidden "Trash" state.

Key Feature

Soft Delete can retain deleted data for a configured period.

Best For

Accidental deletes—you can restore deleted files during the retention period.

Example: Recently Deleted Folder on Your Phone

You are clearing out old photos on your smartphone and accidentally delete a family picture.

How Soft Delete helps:

The photo does not disappear permanently immediately. It stays in the "Recently Deleted" area for the configured period, allowing you to restore it.

10.2 Object Versioning – Your "File History"**What is it?**

Keeps older historical versions of files whenever you overwrite them.

Key Feature

Each version gets its own unique generation number, allowing you to go back to previous versions.

Best For

Accidental overwrites or reverting corrupted master files.

Example: Version History in Microsoft Word

You are writing a project report in Word. You rewrite a whole chapter, save it, and close it.

Later, your boss says:

"Actually, the old version of that chapter was better!"

How Versioning helps:

Because Versioning tracked the previous version, you can select the older version and restore it instead of recreating the content.

10.3 Retention Policies & Locks – Your "Vault"**What is it?**

A strict rule that prevents files from being modified or deleted for a set duration, such as 1 year.

Key Feature

Once locked, files cannot be deleted until the retention period has expired.

Best For

- Legal compliance
- Tax archives
- Ransomware defense

Example: A Time-Capsule Safe with a Digital Timer

You put your company's official tax return files inside a heavy steel safe and set a digital lock timer for 3 years.

How Bucket Lock helps:

The documents cannot be destroyed until the retention period has finished.

10.4 Dual-Region / Multi-Region – Your "Offsite Backup"**What is it?**

Automatically stores copies of your bucket across multiple geographic locations.

Key Feature

If an entire cloud datacenter becomes unavailable, data can remain accessible from another location.

Best For

Critical production systems that require high availability.

Example: Keeping Spare House Keys at a Friend's House

You keep a set of house keys at your home and another set with a friend who lives in another town.

How Dual-Region helps:

If a severe event affects your local location, the copy in another location remains available.

Quick Memory Trick

- **Soft Delete** = Recovers accidentally deleted files.
- **Versioning** = Recovers accidentally overwritten files.
- **Retention Lock** = Stops files from being tampered with.
- **Dual-Region** = Protects against datacenter outages.

11. Google Cloud Storage (GCS) Pricing Model

Google Cloud Storage pricing mainly depends on four factors:

1. **Storage:** Cost for storing data in GCS.
2. **Operations:** Cost for actions such as uploading, reading, copying, and listing objects.
3. **Data Retrieval:** Additional charges for accessing data from colder storage classes.
4. **Network Transfer:** Charges for moving data out of GCS or between regions.

1. Storage Charges

You are charged based on the amount of data stored and the selected storage class.

Storage Class	Approximate Storage Cost	Minimum Storage Duration
Standard	Varies*	None
Nearline	Varies*	30 days
Coldline	Varies*	90 days
Archive	Varies*	365 days

Example:

For 1,000 GB, use the current regional Standard Storage rate to calculate the monthly storage cost. *Cloud Storage pricing varies by location, storage class, and usage.

2. Operations Charges

GCS charges for certain operations performed on buckets and objects.

- **Class A:** Upload, copy, list, and similar operations.
- **Class B:** Read and download operations.
- **Delete:** Generally does not have an operation charge.

The cost depends on the number and type of operations performed.

3. Data Retrieval Charges

Retrieval charges may apply when accessing data stored in colder storage classes.

- **Standard:** No retrieval charge.
- **Nearline:** Retrieval charge applies.
- **Coldline:** Higher retrieval charge.
- **Archive:** Higher retrieval charge.

Therefore, colder storage classes are suitable for data that is accessed less frequently.

4. Network Transfer Charges

Network charges depend on where the data is transferred.

- **Data uploaded to GCS:** Generally free.
- **Data transferred within the same region:** Often free or lower cost, depending on the service.
- **Data transferred between regions:** Charges may apply.
- **Data downloaded to the internet:** Egress charges apply.

Simple Example

For **1 TB (1,000 GB)** of data stored for one month:

Storage Class	Storage Cost	Retrieval Cost	Approximate Total*
Standard	\$20	\$0	\$20
Nearline	\$10	Depends on usage	\$10+
Coldline	\$4	Depends on usage	\$4+
Archive	\$1.20	Depends on usage	\$1.20+

Note: These are simplified example prices. Actual GCS costs vary based on region, storage class, data volume, operations, retrieval, network usage, and applicable pricing rules.

12. Pros and Cons of Google Cloud Storage

Pros

- **Scalable:** Can store very large amounts of data.
- **Highly Durable:** Designed to protect data against loss.
- **Secure:** Supports IAM, encryption, and access controls.
- **Multiple Storage Classes:** Standard, Nearline, Coldline, and Archive for different access patterns.
- **Cost-Effective:** Pay for the storage and services you use.
- **Easy Integration:** Works well with BigQuery, Dataflow, Dataproc, Cloud Functions, and other GCP services.
- **High Availability:** Supports regional, dual-region, and multi-region storage options.
- **Easy to Use:** Data can be managed through the Google Cloud Console, CLI, or APIs.

Cons

- **Network Charges:** Downloading or transferring large amounts of data can increase costs.
- **Retrieval Charges:** Nearline, Coldline, and Archive may have retrieval fees.
- **Minimum Storage Duration:** Some storage classes have minimum storage duration requirements.
- **Not a Traditional File System:** GCS uses object storage rather than a traditional hierarchical file system.
- **Access Configuration:** Incorrect IAM or public-access settings can expose data.
- **Operational Costs:** Large numbers of API operations can increase the overall bill.

13. Steps-by-Steps How to Create the GCS Bucket:

Google Cloud Console (Web UI)

1. Navigate to Cloud Storage:

Log in to the [Google Cloud Console](#), select your target project, and open the navigation menu. Go to Cloud Storage → Buckets, then click Create.

2.Name your bucket: Must be globally unique across all of Google Cloud.

Enter a name conforming to GCS standards (only lowercase letters, numbers, hyphens, and underscores; 3–63 characters; no sensitive info).

3.Choose location type and region:

Select where to store your data:

- **Region:** Lowest latency and cost; best for colocation with BigQuery/Dataflow (e.g., us-central1, asia-south1).
- **Dual-region:** High availability across two specific regions with automatic failover.
- **Multi-region:** Continental availability (e.g., us, eu, asia) for global asset delivery.

4.Select a default storage class:

Choose the default class applied to new uploads:

- **Standard:** Active analytics and frequently read assets (\$0 retrieval fee).
- **Nearline / Coldline / Archive:** Infrequent access, backups, and long-term compliance data.
- Alternatively, check Autoclass to let Google dynamically transition objects between classes.

5.Configure access control and security:

- **Access Control:** Select Uniform (enforces IAM policies across all objects; recommended) instead of Fine-grained (ACLs).
- **Public Access:** Check Enforce public access prevention on this bucket unless the data must be publicly accessible on the internet.

6.Configure data protection and encryption:

- **Data Protection:** Enable Object Versioning or keep Soft Delete enabled (default: 7 days) to protect against accidental deletes.
- **Encryption:** Keep Google-managed encryption key (default) or select a Customer-managed encryption key (CMEK) via Cloud KMS.
- Click Create.

14.Conclusion:

Google Cloud Storage (GCS) is a scalable, secure, and reliable object storage service used to store and manage large amounts of data. It supports different storage classes, flexible access options, and integrates easily with other Google Cloud services. GCS is widely used for data lakes, backups, archives, and data engineering pipelines.

Production-Ready Completion Guide

This addendum follows the same progression used in the Managed Airflow guide: introduction → architecture → core concepts → UI/setup → configuration → security → monitoring → troubleshooting → performance/cost → reliability → production design → end-to-end scenario → common mistakes → operational checklist.

A. Standard learning path

Stage	Expected outcome
1. Introduction	Explain the service, problem and role.
2. Architecture	Explain components and flow.
3. Core concepts	Explain hierarchy, execution and key terminology.
4. UI and setup	Create resources and recognize important settings.
5. Configuration	Configure runtime, identity, networking and scaling.
6. Security	Apply IAM, encryption, secrets and least privilege.
7. Monitoring	Use logs, metrics and alerts.
8. Troubleshooting	Move from symptom to evidence to root cause.
9. Performance and cost	Relate design to throughput, latency and cost.
10. Reliability	Handle retries, duplicates and recovery.
11. Production design	Apply environment separation, CI/CD and ownership.
12. End-to-end scenario	Explain the service in a realistic pipeline.

15. Data lifecycle and protection

Cloud Storage production design should consider lifecycle automation, accidental deletion protection, retention requirements and object versioning.

- Use Object Lifecycle Management or Autoclass where appropriate.
- Understand soft delete, retention policies and holds before deleting data.
- Use versioning when the recovery model requires previous object versions.

16. Access and security design

Bucket access should be designed around IAM, public-access prevention and least privilege.

- Prefer Uniform Bucket-Level Access where it fits the organization.
- Use Public Access Prevention unless public access is intentional.
- Use service accounts or groups instead of individual credentials for applications.

17. Monitoring and operational hygiene

Storage operations should be observable and governed.

- Monitor storage growth, object counts and unusual access patterns.
- Use lifecycle rules to control temporary/intermediate data.
- Review audit logs and access patterns for sensitive buckets.

18. Cost and performance

Storage cost is affected by class, location, operations, retrieval and network transfer.

- Choose storage class from actual access patterns.
- Avoid unnecessary cross-region or internet egress.

- Review operation-heavy workloads such as many small objects.
- Use current Google Cloud pricing for estimates rather than fixed universal rates.

19. Production data governance

Treat buckets as governed data resources, not just file folders.

- Define ownership, retention and deletion policies.
- Document data classification and public/private status.
- Separate raw, curated, temporary and backup data where useful.

20. Common production mistakes

- Publicly exposing a bucket unintentionally.
- Using broad IAM roles.
- Ignoring lifecycle and retention behavior.
- Choosing storage classes only by monthly GB price.
- Ignoring network egress and retrieval charges.
- Deleting data without understanding soft delete/retention/versioning.

21. Operational checklist

- Correct project and bucket selected.
- Location matches latency/residency requirements.
- Storage class matches access pattern.
- IAM/UBLA/public-access settings reviewed.
- Encryption requirements documented.
- Lifecycle/retention policies documented.
- Soft delete/versioning behavior understood.
- Audit/monitoring requirements defined.
- Cost and egress assumptions reviewed.

Official references

<https://docs.cloud.google.com/storage/docs/introduction>

<https://docs.cloud.google.com/storage/docs/access-control/iam>

<https://docs.cloud.google.com/storage/docs/storage-classes>

<https://docs.cloud.google.com/storage/docs/control-data-lifecycles>

Final learning outcome

After completing this document, the learner should be able to explain Google Cloud Storage, create or configure its main resources, operate it from the Google Cloud console, troubleshoot common failures, and describe production controls.